

MOSBAT AI · COMMAND HQ

Your AI Workforce, Explained Simply

A plain-English guide to what the system is, what each capability does, and how it all works together.

Prepared for the Chairman · Positive Nation Economy
Document version 1.0

What's inside

1. The big idea — what Mosbat AI Command HQ is
2. How it works — a company of AI employees
3. The Agents — your AI workforce
4. Skills — the knowledge you give an agent
5. Tools — what agents can actually *do* (each one explained)
6. Memory — how agents remember
7. The Daily Heartbeat — automatic reports
8. What we recently improved
9. How to start the system
10. Glossary — every term in plain words

This guide avoids jargon on purpose. Wherever a technical idea appears, you'll see a 💡
In plain terms note that explains it with an everyday comparison.

1. The big idea

Mosbat AI Command HQ is a **control center for a team of AI assistants** that work for the Positive Nation economy.

Think of it like an online company. Instead of human employees, you have **AI agents** — each one given a job title, a personality, and a set of skills. You (the Chairman) and your managers can chat with any agent, give it work, and watch what it does, all from one dashboard.

Each agent can hold a real conversation, remember what was said before, and — this is the important part — **take real actions in the world**: send an email, search the web, create a document, or inspect your live website.

💡 **In plain terms:** *It's like running a company where you can hire, instruct, and manage a whole staff from a single chat screen — except the staff are AI, available 24/7, and never forget.*

2. How it works — a company of AI employees

When you send a message to an agent, here is what quietly happens behind the scenes, step by step:

#	What happens	The everyday version
1	The system finds the right agent and loads its job description and skills.	The right employee picks up the phone, having read their job manual.
2	It loads the recent conversation history (its memory).	They remember your past chats.
3	It asks the AI brain (Claude) for a response.	The employee thinks and replies.
4	If the reply asks to <i>do</i> something, the system runs that action for real.	They don't just say "I'll email them" — they actually send it.
5	The real result is fed back so the agent reports exactly what happened.	They confirm "Done — here's what I found."
6	The whole exchange is saved to memory.	It's filed away for next time.

 **In plain terms:** *The "AI brain" is a service from Anthropic called Claude. Your system is the office, the desks, the filing cabinet, and the phone lines around that brain.*

3. The Agents — your AI workforce

An **agent** is a single AI employee. Your company already has dozens of them, arranged like a real org chart — a CEO at the top, directors under it, and specialists under them (social media managers, a CFO, QA testers, web developers, and more).

Every agent has:

- **A name and role** — e.g. "Master Brain, Master Control AI" or "QA Quality Specialist".
- **A job description** (called a *system prompt*) — the instructions that shape how it behaves.
- **A manager** — who it reports to, which builds the org chart.
- **A brain setting** — which version of the AI it uses. Smarter, pricier brains for leaders; faster, cheaper ones for routine specialists.

Three brain types in use

Brain	Best for	Think of it as
Opus	The most demanding thinking (your top leaders)	The senior executive
Sonnet	The everyday balance of smart + fast (most staff)	The reliable manager
Haiku	Quick, simple, high-volume tasks	The speedy assistant

You can **hire** new agents, **retire** ones you don't need, edit their instructions, and assign them skills — all from the dashboard.

4. Skills — the knowledge you give an agent

A **skill** is a packet of knowledge or instructions you attach to an agent to make it an expert at something specific.

💡 **In plain terms:** *A skill is like handing an employee a training binder. The agent is generally capable; the skill makes it specialized — "here's exactly how we do social media," or "here are our brand rules."*

How skills work

- You create a skill once (a name + the knowledge text), in the settings panel.
- You assign it to one or more agents.
- From then on, that knowledge is automatically included whenever the agent works — without you re-typing it every time.
- Skills can be turned on or off, edited, or removed at any time.

Skills vs. Tools — don't mix these up:

- A **Skill** is *knowledge* the agent knows (the training binder).
- A **Tool** is an *action* the agent can take (the next section).

5. Tools — what agents can actually do

Most chatbots only talk. Your agents can **act**. A *tool* is a real action an agent can perform. When an agent decides to use one, the system carries it out for real and shows you the result. Here is every tool, in plain language.

Check Site — Live Website QA NEW

The agent visits your real website and reports its true health: Is each page loading? Is it fast? Is the security padlock (SSL) valid? Do the contact forms exist? Are there any broken "page not found" (404) links?

With one extra word (`scope: site`) it reads your site's index and checks **every single page** automatically.

 **In plain terms:** *It's a building inspector that walks through every room of your website and hands back a pass/fail checklist — not a guess, but what it actually saw.*

Web Search

The agent looks things up on the live internet — current facts, prices, news, or contact details — instead of relying only on what it already knew.

 **In plain terms:** *Like an assistant who Googles something for you before answering. (Note: this needs a search key added to turn on — see "What we recently improved.")*

Send Email

The agent sends a real email — a report, a notification, or outreach — to any address you specify.

 **In plain terms:** *It doesn't just say "I'll email them." It actually sends it from your company mailbox.*

Generate Image

The agent creates a brand-new picture from a written description — useful for social posts, thumbnails, or mock-ups.

 **In plain terms:** *Like a graphic designer who sketches exactly what you describe, on demand.*

Generate Document

The agent writes a formatted document — a report, proposal, or plan — and saves it as a file you can open and print to PDF.

 **In plain terms:** *Like dictating to an assistant who hands you back a clean, ready-to-print document.*

Create Task

An agent (often a manager) hands a job to another agent. The work goes into a queue and gets done.

 **In plain terms:** *A manager delegating to a team member — "you handle this, report back when done."*

Hire Agent / **Fire (Deactivate) Agent**

Agents (or you) can expand the team by creating a new agent, or deactivate one that's no longer needed. Deactivation is reversible.

 **In plain terms:** *Growing or trimming your staff — except retiring someone can be undone with a click.*

6. Memory — how agents remember

Each agent keeps an ongoing memory of your conversations, so it doesn't start from scratch every time.

It works on two levels:

- **Recent memory** — the last several messages are always kept in full, word for word.
- **Long-term memory** — when a conversation gets long, older parts are automatically condensed into a short summary so the gist is never lost, but it stays affordable.

 **In plain terms:** *Like an assistant who remembers your recent chats exactly, and keeps brief notes on everything older — so they always have context without re-reading a year of transcripts.*

Memory now stays consistent across days and devices — chat from your phone today and your laptop tomorrow, and the agent picks up where you left off.

7. The Daily Heartbeat — automatic reports

The system can wake an agent on a schedule (set for 5:00 PM) to produce a daily report on its own, without anyone asking.

💡 **In plain terms:** *Like a manager who emails you an end-of-day summary every evening — automatically.*

Where this is headed: today's heartbeat *writes* a summary. The natural next step is to combine it with the new **Check Site** tool so every morning you receive a *real, measured* status report on the live website — automatically. This is ready to switch on whenever you like.

8. What we recently improved

Recent work made the system actually work, cost less, and tell the truth. In plain language:

① We brought the agents back online

The agents were completely silent. Two problems were stacked: the access key to the AI brain was invalid, and one of the top agents used a newer brain that rejected an old setting. Both were fixed. **All agents now reply.**

② We made it cheaper and faster

Each agent carries a large set of instructions that used to be re-sent on every single message — like re-faxing a 10-page cover sheet before every sentence. We turned on **reuse**, so those instructions are remembered for a few minutes and served back at about a tenth of the cost.

💡 **In plain terms:** *We verified it works: the second message reused the instructions instead of paying full price again.*

③ We gave agents real eyes on the website

Before, an agent asked to "check the site" could only guess. Now it genuinely visits the live site and reports verifiable facts. We tested it against positivenation.org: it correctly found all 40 pages and confirmed they load with no broken links.

④ We made the reports trustworthy

Previously an agent might describe results before it had them — essentially predicting. Now the **real** findings are fed back to the agent before it writes its report, so the numbers you read are the numbers that were actually measured.

⑤ We removed false alarms

The first website check raised a scary "broken form / lost leads" alert. It was wrong — your site is built with Wix, which behaves in ways that fooled a basic checker. We tested against your live site, found three such quirks, and taught the tool to recognize them, so it no longer cries wolf.

The false alarm	The real explanation
"Form is broken (404)"	The form submits via modern website code — perfectly normal.
"All pages are broken"	Wix answers a quick "knock" with an error but a real "visit" with success — we switched to a real visit.
"No pages found"	The address with and without "www" was treated as two different sites — now unified.

9. How to start the system

This is a website application. It is **not** started with the Node/ `pnpm` commands you may have seen elsewhere — those are for a different kind of project.

To run it

- 1. Open a terminal in the project folder.
- 2. Type: `php spark serve`
- 3. Open `http://localhost:8080` in your browser.

Other handy commands

Command	What it does
<code>php spark serve</code>	Start the dashboard and chat.
<code>php spark heartbeat:run</code>	Generate the daily report by hand.
<code>php spark list</code>	Show all available commands.

10. Glossary — every term in plain words

Term	Plain meaning
Agent	One AI employee with a name, a role, and a job description.
Skill	A packet of knowledge you attach to an agent to make it a specialist.
Tool	A real action an agent can take (email, search, website check, etc.).
System prompt	An agent's job description — the instructions that shape its behavior.
Claude	The outside AI "brain" (from Anthropic) that powers the agents' thinking.
Memory	The saved history of your conversations with an agent.
Heartbeat	A scheduled, automatic check-in that produces a report.
Dashboard	The web screen where you see the org chart and chat with agents.
404	A "page not found" error — a broken or missing web page.
SSL	The security padlock that proves a website is encrypted and safe.
Prompt caching	Reusing an agent's instructions for a short time to save money and time.
Supabase	The online database where agents, skills, and memory are stored.